

Curve, Freeze, Thaw: HPO with FT-PFN

Yahya Fati Haji

August 2, 2026

Problem Formulation

Let Λ be a hyperparameter search space and $f(\lambda, b)$ the performance (validation accuracy, in this project) of configuration $\lambda \in \Lambda$ after being trained for b discrete training steps (epochs, here). A configuration's *learning curve* is the sequence $\{f(\lambda, 1), f(\lambda, 2), \dots\}$ as b grows.

Freeze-thaw HPO relaxes the usual "pick a config, train it to completion, observe one number" protocol. Instead, resources are spent **incrementally**: at every iteration the optimizer either **thaws** (resumes/continues) a previously partially-trained, currently-frozen configuration for one more step, or starts a brand-new configuration.

Formally, the goal is to find a resource allocation that maximizes

$$\max_{\lambda \in \Lambda, 1 \leq b \leq b_\lambda} f(\lambda, b)$$

subject to

$$N = |\{\text{decisions made}\}| \leq B, \quad b_\lambda^{\min} \leq b_\lambda \leq b_\lambda^{\max}, \\ \forall \lambda \in \Lambda \text{ with } b_\lambda > 0, \quad b_\lambda \in \mathbb{Z}_{\geq 0}, \quad \forall \lambda \in \Lambda$$

- a total trial budget, B - the number of freeze/thaw decisions the optimizer is allowed to make.
- $b_\lambda^{\min}$ and $b_\lambda^{\max}$ - per-configuration epoch bounds that prevent premature judgments on too little training and wasted spend on configs that have already converged.

The FT-PFN Surrogate

Rather than fitting a Gaussian Process or random forest online (as SMAC/BOHB do), ifBO uses **FT-PFN**, a *Prior-data Fitted Network*: a transformer trained once, offline, on large quantities of **synthetic** learning-curve data sampled from a hand-designed prior, and then used purely for inference at HPO time via *in-context learning* — no weights are updated during the actual HPO run.

The implementation follows **ifBO** as introduced in:

H. Rakotoarison, S. Adriaensen, N. Mallik, S. Garibov, E. Bergman, F. Hutter. *In-Context Freeze-Thaw Bayesian Optimization for Hyperparameter Optimization*. ICML 2024. <https://arxiv.org/abs/2404.16795>

What the surrogate models

FT-PFN approximates the posterior predictive distribution

$$p(f(\lambda_{\text{test}}, b_{\text{test}}) \mid \lambda_{\text{test}}, b_{\text{test}}, H) \tag{1}$$

i.e., "given everything observed so far about *other* (and this) configurations' partial learning curves, what is the distribution over this configuration's performance at some future training step b_{test} ?" The set H is passed to the network as a sequence of tokens (its **context**) at inference time; the transformer's attention mechanism implicitly performs Bayesian updating over this context in a single forward pass, rather than through iterative model refitting. This is what makes each acquisition query cheap enough to run every freeze-thaw step (the paper reports 10–100× speedups over refitting-based gray-box surrogates such as DPL and DyHPO).

Acquisition function: MFPI and MFPI-random

The paper’s definition

Multi-fidelity Probability of Improvement, as defined in the paper (Eq. 3):

$$\text{MFPI}(\lambda; h, T) = P(M(\lambda, \min(b_\lambda + h, b_{\max})) > T) \quad (2)$$

— the surrogate-predicted probability that configuration λ , if thawed for h more steps, would exceed a target performance T . Two “hyper-hyperparameters” govern it: the lookahead horizon h and the improvement target T . Rather than fixing these, the paper proposes **MFPI-random** (Eq. 4), redrawing both **every single freeze-thaw iteration**:

$$\text{MFPI-random}(\lambda) = \text{MFPI}(\lambda; h_{\text{rand}}, T_{\text{rand}}) \quad (3)$$

$$h_{\text{rand}} \sim U(1, b_{\max}) \quad (4)$$

$$T_{\text{rand}} = f_{\text{best}} + \tau_{\text{rand}} \cdot (1 - f_{\text{best}}) \quad (5)$$

$$\log_{10}(\tau_{\text{rand}}) \sim U(-4, -1) \quad (6)$$

where f_{best} is the best performance observed so far. This amounts to sampling an acquisition function from an implicit *portfolio* of MFPI instances at every step, which the paper’s ablations (Figure 4) show is necessary — fixed-horizon or fixed-threshold variants, and standard Expected Improvement paired with FT-PFN’s heavy-tailed posteriors, both underperform substantially.

Growing the candidate pool

The paper’s Algorithm 1 (see the pseudocode in the appendix of the original paper) is stated over the *whole* search space Λ implicitly — at each iteration the acquisition is (conceptually) maximized over all $\lambda \in \Lambda$, which naturally lets brand-new, never-tried configurations compete against partially-trained ones for being selected next.

This implementation instead maintains an **explicit, dynamically growing list** of candidates (starting empty), and layers an **explicit decaying ϵ -greedy exploration floor** on top of MFPI-random rather than substituting for it:

$$\epsilon(t) = \epsilon_{\min} + (\epsilon_0 - \epsilon_{\min}) \cdot \left(1 - \frac{t}{T}\right)^p \quad (7)$$

$$\text{where } \epsilon_{\min} = 0.05, \quad p = 2, \quad T = N$$

ϵ_0 (ifbo_initial_epsilon, default 1.0) means the very first steps are pure exploration — with an empty/small pool there’s nothing meaningful yet for MFPI-random to discriminate between — decaying polynomially toward a floor of 0.05 so some unconditional exploration always remains, even late in the run. Each iteration then branches:

- **With probability $\epsilon(t)$** (the exploration floor): a brand-new configuration is sampled uniformly from Λ , added to the pool with zero observations, and selected directly — bypassing the surrogate entirely.
- **With probability $1 - \epsilon(t)$** (an exploitation round): the *pending* pool (candidates not yet at $b_{\max}$, minus any already claimed earlier in the same parallel batch, and minus the current best-so-far candidate once it meets the incumbent-exclusion rule below) is assembled. MFPI-random ($h_{\text{rand}}, T_{\text{rand}}$ redrawn as above) then scores every member of that pool against the FT-PFN context in a single batched query, and takes the arg max over the PI scores (or a softmax draw over them). If the pending pool is empty this round (everything left is either maxed out or the just-excluded incumbent), a fresh configuration is sampled as a fallback instead, the same as an exploration step.

One more rule guards the exploitation branch: once the current best-so-far candidate has accumulated at least ifbo_incumbent_exclusion_min_observations (default 2) freeze-thaw steps, it is temporarily dropped from the pending pool for that round. Left unchecked, it tends to keep re-winning $\text{PI}(T_{\text{rand}})$ against its own already-confirmed best (since T_{rand} sits just above f_{best}), sinking budget into repeatedly re-thawing itself instead of advancing or discovering other candidates.

A fresh candidate no longer competes in exploitation rounds — neither of these two rules is part of the published ifBO method, which is stated over the whole search space Λ implicitly rather than an

explicit, growing candidate list. An earlier version of this implementation *did* throw one freshly-sampled, not-yet-pooled configuration into every exploitation round’s competition alongside **pending**, reasoning that it would let exploitation rounds surface genuinely new regions of the space too, without unconditionally growing the pool (and the FT-PFN context) on every round the way pure ϵ -driven injection would. In practice this backfired specifically under greedy selection, and a seed sweep in **sample-results/seeded/** (**ifbo-greedy** vs. **ifbo-random**, matched seeds on **amazon**) traced why: once the incumbent-exclusion rule above walls off the current best, every *pending* candidate is a known quantity the surrogate confidently ranks below the now-frozen target — but a never-observed candidate’s predictive distribution under FT-PFN is wide/uncertain, which inflated its PI score against that same high threshold relative to an already-observed, confidently-mediocre pending candidate.

Search space: the sequence-dl approach

sequence-dl is a BiLSTM-with-attention text classifier trained **from scratch** — no fine-tuning of a pretrained encoder — so its search space has to cover both generic optimization knobs and the architecture of a recurrent encoder built from nothing. $\Lambda_{\text{sequence-dl}}$ splits into two groups: hyperparameters shared with **transformer** and hyperparameters specific to the BiLSTM.

Shared hyperparameters (both approaches sample these; ranges below are shared, defaults only sometimes differ per-approach):

dropout

Range: [0.0, 0.5] | **Default:** 0.2

Rationale: Regularizes the classifier head / recurrent stack.

weight_decay

Range: $[10^{-6}, 10^{-2}]$ (log scale) | **Default:** 10^{-4}

Rationale: Standard L2 regularization;

optimizer

Range: {adam, adamw, sgd} | **Default:** adamw

Rationale: —

scheduler

Range: {stepplr, cosine, exponentiallr, reducelronplateau} | **Default:** cosineannealinglr

Rationale: Algorithm Selection

batch_size

Range: [32, 512] (log scale) | **Default:** 64

Rationale: —

max_seq_length

Range: [64, 256] (log scale) | **Default:** 128

Rationale: Bounds both compute (LSTM cost is linear in sequence length after packing) and how much of a long document is even visible to the model; log-scaled so short/long regimes get comparable sampling density.

warmup_ratio

Range: [0.0, 0.2] | **Default:** 0.1

Rationale: Fraction of total steps spent on linear LR warmup before the main schedule kicks in — guards against the early, high-variance gradients typical of a randomly-initialized model.

sequence-dl-Specific Hyperparameters

hidden_dim

Range: [32, 256] (log scale) | **Default:** 128

Rationale: LSTM hidden size per direction (the model is always bidirectional, so the pooled representation is $2\times$ this).

learning_rate

Range: $[10^{-4}, 10^{-2}]$ (log scale) | **Default:** 10^{-3}

Rationale: —

seq_embed_dim

Range: [32, 512] (log scale) | **Default:** 128

Rationale: Token embedding dimensionality. Independent of any pretrained encoder’s native hidden size

on purpose — see `seq_pretrained_model_name` below for how a mismatch against the source embedding matrix is handled.

`seq_num_layers`

Range: [1, 3] | **Default:** 1

Rationale: Stacked LSTM layers. Kept shallow (max 3) because a from-scratch recurrent stack is harder to optimize as it deepens (vanishing gradients through both time and depth), and because HPO wall-clock is a hard constraint — depth is one of the more expensive ways to spend that budget for the accuracy it typically buys on these datasets.

`seq_pretrained_model_name`

Range: {distilbert-base-uncased, bert-base-uncased, google/bert_uncased_L-4_H-512_A-8, microsoft/xtrem

Default: distilbert-base-uncased

Rationale: Does **not** select a fine-tuned backbone (`sequence-dl` never runs one) — it picks which pre-trained model’s **WordPiece tokenizer and token-embedding matrix** warm-start the from-scratch BiLSTM. The tokenizer and embedding source are always the same model, since the BiLSTM’s vocab indices must line up with whichever embedding matrix seeds it. When the chosen model’s native embedding dimensionality doesn’t match the sampled `seq_embed_dim`, the embedding matrix is PCA/SVD-projected down (or randomly padded up) to fit — see `_pretrained_embedding_init` in `sequence_dl.py` — which preserves the directions of highest variance in the pretrained embedding space instead of discarding the warm start entirely.

The diagram below shows `BiLSTMClassifier`’s data flow (`automl/core/approaches/sequence_dl.py`): a warm-started embedding, a bidirectional LSTM sized by `hidden_dim/seq_num_layers`, additive (Bahdanau-style) attention pooling over every timestep rather than just the final hidden state, then dropout and a linear classifier.



Figure 1: `BiLSTMClassifier` architecture

Two omissions are deliberate, not oversights: `max_grad_norm` (gradient-clipping threshold) is fixed by CLI/-config default rather than tuned — it’s a stability safeguard, not an accuracy lever, so putting it in Λ would spend trials on a dimension unlikely to move the metric. And unlike `transformer`, there is no `freeze_ratio`-equivalent knob: nothing is pretrained end-to-end here to freeze in the first place, only the embedding matrix, and that is warm-started rather than frozen.

Code Flow

Putting the pieces from the previous sections together, this is what a full search actually does end to end, from an empty candidate pool to a finished submission.

The search starts with nothing: no configuration has been tried yet, and the candidate pool is empty. Every iteration opens with a single weighted coin flip that decides whether that step explores or exploits — the odds start out heavily favoring exploration, since with an empty pool there is nothing yet worth exploiting, and gradually shift toward exploitation as the run progresses and a real pool of partially-trained candidates builds up.

An **exploration** step is the simple branch: a brand-new configuration is drawn at random from the search space, dropped straight into the pool, and trained for one step immediately — no model, no scoring, no comparison against anything else.

An **exploitation** step is where the surrogate does its work. Every candidate already in the pool that hasn’t yet been trained to the maximum allowed budget is treated as a contender — except the current best-so-far candidate, once it’s accumulated enough freeze-thaw steps to trigger the incumbent-exclusion rule (§Growing the candidate pool). All of them are laid before the surrogate model at once, alongside the complete history of every partial learning curve observed anywhere in the search so far. The surrogate — a model trained once, in advance, and never updated again during the search itself — reads that history and, for each contender, predicts how likely it is to cross a target performance level within some number of additional training steps. Neither the target nor the horizon is fixed: both are redrawn at random on every single round, so the search never commits to one fixed idea of “how much better” or “how far ahead,” and instead samples from a broad spread of possible answers each time. Whichever contender comes out on top is the one trained further. If there are no contenders at all this round (every pending candidate is either maxed out or the just-excluded incumbent), a brand-new configuration is sampled instead, the same as an exploration step.




Figure 2: The ifBO search loop

Whichever branch fires, the chosen configuration is then trained for exactly one more increment, its progress is checkpointed, and the new point on its learning curve is appended to the shared history the surrogate will read on the next round. This one-step-at-a-time cycle — decide, maybe consult the surrogate, train a little further, record the result — repeats until the overall budget of decisions runs out.

At the end of the run, rather than keeping only the single best configuration ever seen, the strongest handful found across the whole search are retrained to completion and combined by majority vote, and it is that ensemble’s predictions that make up the final submission.

Results

Budget:

```

max_budget: 10
min_budget: 3
n_trials: 20

```

Testbed Results